

FE590. Assignment #2.

2021-10-23

I pledge on my honor that I have not given or received any unauthorized assistance on this assignment/examination. I further pledge that I have not copied any material from a book, article, the Internet or any other source except where I have expressly cited the source.

By filling out the following fields, you are signing this pledge. No assignment will get credit without being pledged.

Name: Muhammet Furkan Isik

CWID: 10472193

Date: 10/21/2021

Instructions

In this assignment, you should use R markdown to answer the questions below. Simply type your R code into embedded chunks as shown above. When you have completed the assignment, knit the document into a PDF file, and upload both the .pdf and .Rmd files to Canvas.

```
CWID = 10472193 #Place here your Campus wide ID number, this will personalize  
#your results, but still maintain the reproduceable nature of using seeds.  
#If you ever need to reset the seed in this assignment, use this as your seed  
#Papers that use -1 as this CWID variable will earn 0's so make sure you  
change  
#this value before you submit your work.  
personal = CWID %% 10000  
set.seed(personal) #You can reset the seed at any time in your code,  
#but please always set it to this seed.
```

Question 1

You are part of a research team of a FinTech company dedicated to forecast the price direction of public US companies. Using the attached dataset, select the company assigned to you according to your Stevens ID in the spreadsheet "assignments".

You must build a forecasting model of your assigned company's next day's excess return (stock return – risk free rate; ExRet) and its direction using the years 2014-2018 to train and 2019 to test your model.

The dataset is based on the Fama French 3 factor model:

$$R_{it} - R_{ft} = \alpha_{it} + \beta_1(R_{Mt} - R_{ft}) + \beta_2SMB_t + \beta_3HML_t + \epsilon_{it}$$

where: R_{it} =total return of a stock or portfolio i at time t R_{ft} =risk free rate of return at time t
 RM_t =total market portfolio return at time t $R_{it}-R_{ft}$ =expected excess return
 RM_t-R_{ft} =excess return on the market portfolio (index) SMB_t =size premium (small minus big) HML_t =value premium (high minus low) $\beta_{1,2,3}$ =factor coefficients

The dataset includes the following variables or columns: Ticker: identifier of the stock
 Permno: permanent number of the stock Date: date of observation Alpha: Alpha of stock i
 b_{mkt} : Beta on mkt_{rft} b_{smb} : Beta on SMB_t b_{hml} : Beta on HML_t $ivol$: Idiosyncratic
 Volatility: Variance of residuals or the part that cannot be explained by the Fama French
 model $tvol$: Total Volatility: Variance of stock returns $Exret$: Excess Return from Risk
 Model: stock return ($price_t/price_{t-1}-1$) – risk free rate

For this exercise you do not need to know the details of the Fama French model. However, further details and data can be obtained at:

http://mba.tuck.dartmouth.edu/pages/faculty/ken.french/Data_Library/f-f_factors.html
https://mba.tuck.dartmouth.edu/pages/faculty/ken.french/data_library.html

1. Create a csv file with the information of the company assigned to you according to your Stevens ID in the tab “assignments” of the provided database. Read this csv file in R. Sort the variables by date in a new dataset. List the names of the variables in the dataset.

Importing raw data

```
library(readxl)
row_data <- read_excel("stocks100_2014_19.xlsx")
```

Subsetting only MDT Ticker

```
data= row_data[row_data$TICKER== "MDT", ]
```

Ordering the dataset by date

```
data= data[order(data$DATE), ]
data= data[,c(2,1,3,4,5,6,7,8,9,10)]
```

```
names(data)
```

```
## [1] "DATE" "PERMNO" "alpha" "b_mkt" "b_smb" "b_hml" "ivol" "tvol"
## [9] "exret" "TICKER"
```

```
head(data)
```

```
## # A tibble: 6 x 10
```

```
##       DATE PERMNO  alpha b_mkt  b_smb  b_hml   ivol   tvol   exret
TICKER
##    <dbl>  <dbl>  <dbl> <dbl>  <dbl>  <dbl>  <dbl>  <dbl>  <dbl>
<chr>
```

```
## 1 20140102  60097 0.0004 0.928 -0.234 -0.0669 0.00775 0.00995 0.00500 MDT
## 2 20140103  60097 0.0004 0.928 -0.220 -0.0661 0.00778 0.00998 0.0198  MDT
```

```
## 3 20140106 60097 0.0005 0.924 -0.242 -0.0454 0.00785 0.0100 0.0171 MDT
## 4 20140107 60097 0.0005 0.929 -0.237 -0.0346 0.00781 0.0100 0.00496 MDT
## 5 20140108 60097 0.0005 0.929 -0.238 -0.0425 0.00787 0.0101 0.0163 MDT
## 6 20140109 60097 0.0004 0.928 -0.248 0.0061 0.00802 0.0102 -0.0242 MDT
```

2. Generate a new variable `exret1` which is the excess return of the next day and `exret_sq` which is squared of `exret`. As the variables `PERMNO`, `DATE` and `TICKER` will be unimportant, remove those fields from your data frame.

```
# Removing PERMNO, DATE and TICKER
```

```
data= subset(data, select= - c(DATE, PERMNO, TICKER) )
head(data)
```

```
## # A tibble: 6 x 7
##   alpha b_mkt b_smb b_hml ivol tvol exret
##   <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 0.0004 0.928 -0.234 -0.0669 0.00775 0.00995 0.00500
## 2 0.0004 0.928 -0.220 -0.0661 0.00778 0.00998 0.0198
## 3 0.0005 0.924 -0.242 -0.0454 0.00785 0.0100 0.0171
## 4 0.0005 0.929 -0.237 -0.0346 0.00781 0.0100 0.00496
## 5 0.0005 0.929 -0.238 -0.0425 0.00787 0.0101 0.0163
## 6 0.0004 0.928 -0.248 0.0061 0.00802 0.0102 -0.0242
```

```
# Creating exret1 variable
```

```
e1= data[2:nrow(data), "exret"]
```

```
# Creating exret_sq variable
```

```
e2= e1^2
```

```
# Merging e1 and e2
```

```
e_1_2= cbind(e1,e2)
```

```
colnames(e_1_2)= c("exret1", "exret_sq")
```

```
# Adding them into the dataframe
```

```
data= data[1:nrow(data)-1, ]
```

```
data= cbind(data, e_1_2)
```

```
head(data)
```

```
##   alpha b_mkt b_smb b_hml ivol tvol exret exret1
## 1 4e-04 0.9282 -0.2337 -0.0669 0.007747 0.009945 0.005005 0.019830
## 2 4e-04 0.9276 -0.2201 -0.0661 0.007783 0.009977 0.019830 0.017146
## 3 5e-04 0.9239 -0.2423 -0.0454 0.007853 0.010012 0.017146 0.004955
```

```
## 4 5e-04 0.9287 -0.2368 -0.0346 0.007810 0.010007 0.004955 0.016325
## 5 5e-04 0.9289 -0.2375 -0.0425 0.007874 0.010052 0.016325 -0.024235
## 6 4e-04 0.9279 -0.2480 0.0061 0.008023 0.010178 -0.024235 0.008848
##      exret_sq
## 1 3.932289e-04
## 2 2.939853e-04
## 3 2.455203e-05
## 4 2.665056e-04
## 5 5.873352e-04
## 6 7.828710e-05
```

3. What is the range of each quantitative variable? Answer this question using the range() function with the sapply() function e.g., sapply(cars, range). Print a simple table of the ranges of the variables. The rows should correspond to the variables. The first column should be the lowest value of the corresponding variable, and the second column should be the maximum value of the variable. The columns should be suitably labeled.

```
# Table for range data using sapply
range_data= sapply(data,range)

# transpose of data
range_data= t(range_data)

# Changing colnames
colnames(range_data)= c("Min", "Max")
#range_data

# Converting it into a dataframe
as.data.frame(range_data)

##           Min      Max
## alpha -9.0000e-04 0.000800000
## b_mkt  5.6470e-01 1.113300000
## b_smb -4.5510e-01 0.035000000
## b_hml -6.7520e-01 0.086900000
## ivol  7.6140e-03 0.010500000
## tvol  9.3570e-03 0.013803000
## exret -8.5106e-02 0.056594000
## exret1 -8.5106e-02 0.056594000
## exret_sq 2.5000e-11 0.007243031
```

4. What is the mean and standard deviation of each variable? Create a simple table of the means and standard deviations.

```
# Mean and Sd for each variable
mean_data= sapply(data, mean)
```

```

sd_data= supply(data,sd)
# Taking transpose
t(mean_data)

##           alpha      b_mkt      b_smb      b_hml      ivol      tvol
## [1,] 0.000111332 0.8506675 -0.1876306 -0.3166404 0.009031587 0.01162728
##           exret      exret1      exret_sq
## [1,] 0.0001735394 0.0001723764 8.471263e-05

t(sd_data)

##           alpha      b_mkt      b_smb      b_hml      ivol      tvol
## [1,] 0.0003497974 0.1378713 0.1056678 0.1819759 0.0006319122 0.001058756
##           exret      exret1      exret_sq
## [1,] 0.009205883 0.009205383 0.0003089153

# Creating a dataframe
table= cbind(mean_data,sd_data)
as.data.frame(table)

##           mean_data      sd_data
## alpha      1.113320e-04 0.0003497974
## b_mkt      8.506675e-01 0.1378712588
## b_smb     -1.876306e-01 0.1056677970
## b_hml     -3.166404e-01 0.1819759045
## ivol      9.031587e-03 0.0006319122
## tvol      1.162728e-02 0.0010587562
## exret      1.735394e-04 0.0092058829
## exret1     1.723764e-04 0.0092053830
## exret_sq   8.471263e-05 0.0003089153

```

5. Split your data into a 70% training set and a 30% testing set. Using the regsubsets function in the leaps library, regress next day excess return on the remaining variables only using the training sample.

```

# Splitting the dataset

# Creating training dataset 70%
train= 0.7* nrow(data)
train=round(train)

data_train= data[1: train, ]
tail(data_train)

##           alpha b_mkt  b_smb  b_hml      ivol      tvol      exret      exret1
## 1051 -6e-04 0.8325 -0.1219 -0.2766 0.008830 0.010394 0.006470 0.008209
## 1052 -6e-04 0.8318 -0.1072 -0.2843 0.008847 0.010409 0.008209 0.008676
## 1053 -6e-04 0.8373 -0.1151 -0.2751 0.008844 0.010423 0.008676 -0.002445
## 1054 -6e-04 0.8342 -0.1087 -0.2742 0.008837 0.010438 -0.002445 0.002830
## 1055 -6e-04 0.8336 -0.1056 -0.2750 0.008840 0.010437 0.002830 0.007661

```

```
## 1056 -6e-04 0.8277 -0.1021 -0.2746 0.008854 0.010437 0.007661 -0.010955
##          exret_sq
## 1051 6.738768e-05
## 1052 7.527298e-05
## 1053 5.978025e-06
## 1054 8.008900e-06
## 1055 5.869092e-05
## 1056 1.200120e-04

# Creating test dataset 30%
test= train+1
data_test= data[test: nrow(data),]

tail(data_test)

##      alpha  b_mkt  b_smb  b_hml  ivol  tvol  exret  exret1
## 1504 1e-04 0.6291 -0.4540 -0.3131 0.009103 0.011031 -0.003217 0.000634
## 1505 1e-04 0.6280 -0.4545 -0.3124 0.009102 0.010987 0.000634 0.000433
## 1506 2e-04 0.6043 -0.4185 -0.3261 0.009063 0.010791 0.000433 -0.010935
## 1507 2e-04 0.6032 -0.4074 -0.3288 0.009090 0.010608 -0.010935 0.005024
## 1508 2e-04 0.5978 -0.3991 -0.3294 0.009079 0.010568 0.005024 0.002287
## 1509 2e-04 0.5967 -0.3972 -0.3279 0.009080 0.010566 0.002287 0.003250
##          exret_sq
## 1504 4.019560e-07
## 1505 1.874890e-07
## 1506 1.195742e-04
## 1507 2.524058e-05
## 1508 5.230369e-06
## 1509 1.056250e-05
```

a. Print a table showing what variables would be selected using best subset selection for all predictors of the training set. Determine the optimal model using Mallows' Cp and output the model, including its coefficients.

```
# Best Subset Selection using training dataset
library(leaps)
ex.sub= regsubsets(exret~., data=data_train, nvmax=10)
t(summary(ex.sub)$which)

##           1      2      3      4      5      6      7      8
## (Intercept) TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE
## alpha      FALSE TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE
## b_mkt      FALSE FALSE FALSE FALSE FALSE FALSE TRUE  TRUE
## b_smb      FALSE FALSE FALSE TRUE  TRUE  TRUE  TRUE  TRUE
## b_hml      FALSE FALSE FALSE FALSE FALSE FALSE FALSE TRUE
## ivol       FALSE FALSE FALSE FALSE FALSE TRUE  TRUE  TRUE
## tvol       FALSE FALSE FALSE FALSE TRUE  TRUE  TRUE  TRUE
## exret1     TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE
## exret_sq   FALSE FALSE TRUE  TRUE  TRUE  TRUE  TRUE  TRUE
```

Mallow's C_p



For a fitted least squares model containing d predictors, the C_p estimate of test MSE is computed using the equation:

$$C_p = \frac{1}{n}(RSS + 2d\hat{\sigma}^2)$$

where $\hat{\sigma}^2$ is an estimate of the variance of the error

Mallow's C_p is sometimes defined as

$$C'_p = \frac{RSS}{\hat{\sigma}^2} + 2d - n$$

The model with the smallest C_p will also have the smallest C'_p

Mallow's C_p

- According to the Mallows' C_p Model 3 should be selected, which gives the minimum value

```
## According to the Mallows' Cp Model 3 should be selected
# Then our model is exret~ beta0 + b1*alpha + b2*exret1 + b3* exret_sq
```

```
# Using Mallows' Cp to determine the best fitted model
cp= summary(ex.sub)$cp
cp
```

```
## [1] 6.291117 1.274103 1.072380 2.899079 4.630677 6.199525 7.001561
9.000000
```

```
# Choosing the minimum Mallow's value
which.min(cp)
```

```
## [1] 3
```

```
#bic
#bic=summary(ex.sub)$bic
```

- Then our model is $\text{exret} \sim \text{beta0} + \text{b1} * \text{alpha} + \text{b2} * \text{exret1} + \text{b3} * \text{exret_sq}$
- $\text{exret} = -8.569053\text{e-}05 * \text{beta0} + 2.072522\text{e+}00 * \text{alpha} - 9.578236\text{e-}02 * \text{exret1} - 1.411347\text{e+}00 * \text{exret_sq}$

```
# Then our model is exret~ beta0 + b1*alpha + b2*exret1 + b3* exret_sq
# exret= -8.569053e-05 * beta0 + 2.072522e+00* alpha -9.578236e-02 *exret1
```

```
-1.411347e+00 * exret_sq
```

```
cp_model= lm(exret~ alpha+exret1+ exret_sq, data=data_train)
# summary(cp_model)$coefficients
coef(cp_model)
```

```
## (Intercept)          alpha          exret1          exret_sq
## -8.569053e-05  2.072522e+00 -9.578236e-02 -1.411347e+00
```

b. Print a table showing what variables would be selected using forward subset selection for all predictors of the training set. Determine the optimal model using BIC and output the model, including its coefficients.

Bayesian Information Criterion (BIC)



For the least squares model with d predictors up to irrelevant constants

$$BIC = \frac{1}{n}(RSS + d\hat{\sigma}^2 \log n)$$

Note that BIC replaces the $2d\hat{\sigma}^2$ with $d\hat{\sigma}^2 \log n$.

Since $\log n > 2$ for $n > 7$, the BIC places a heavier penalty on models with many variables

Bayesian Information Criteria

Used regsubset function and specified forward method to implement forward step wise implementation

```
for.sub= regsubsets(exret~., data=data_train, nvmax=10,method = "forward")
t(summary(for.sub)$which)
```

```
##           1      2      3      4      5      6      7      8
## (Intercept) TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE
## alpha      FALSE TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE
## b_mkt      FALSE FALSE FALSE FALSE FALSE FALSE TRUE  TRUE
## b_smb      FALSE FALSE FALSE TRUE  TRUE  TRUE  TRUE  TRUE
## b_hml      FALSE FALSE FALSE FALSE FALSE FALSE FALSE  TRUE
## ivol       FALSE FALSE FALSE FALSE FALSE TRUE  TRUE  TRUE
```



```
## tvol      FALSE FALSE FALSE FALSE  TRUE  TRUE  TRUE  TRUE
## exret1    TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE
## exret_sq  FALSE FALSE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE
```

- Used BIC to come up with the model, according to BIC model with two predictors alpha1 and exret1 and constant term

Used BIC to come up with the model, according to BIC model with two predictors alpha1 and exret1 and constant term

```
bic= summary(for.sub)$bic
which.min(bic)
```

```
## [1] 2
```

- coefficients (Intercept) : -0.0002114784 alpha: 2.1030051398 exret1: -0.0868634902

Model summary and coefficients (Intercept) : -0.0002114784 alpha: 2.1030051398 exret1: -0.0868634902

```
bic_model= lm(exret~alpha+exret1,data=data_train)
summary(bic_model)
```

```
##
```

```
## Call:
```

```
## lm(formula = exret ~ alpha + exret1, data = data_train)
```

```
##
```

```
## Residuals:
```

```
##      Min       1Q   Median       3Q      Max
## -0.082243 -0.004553 -0.000072  0.004535  0.043282
```

```
##
```

```
## Coefficients:
```

```
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -0.0002115  0.0003031  -0.698  0.48544
## alpha       2.1030051  0.7932461   2.651  0.00814 **
## exret1     -0.0868635  0.0305820  -2.840  0.00459 **
```

```
## ---
```

```
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
##
```

```
## Residual standard error: 0.009207 on 1053 degrees of freedom
```

```
## Multiple R-squared:  0.01435,    Adjusted R-squared:  0.01248
```

```
## F-statistic: 7.666 on 2 and 1053 DF,  p-value: 0.0004951
```

```
coef(bic_model)
```

```
##      (Intercept)      alpha      exret1
## -0.0002114784  2.1030051398 -0.0868634902
```

c. Print a table showing what variables would be selected using backward subset selection for all predictors of the training set. Determine the optimal model using adjusted R^2 and output the model, including its coefficients.

Adjusted R^2



Recall:

$$R^2 = 1 - \frac{RSS}{TSS}$$

Note that this will always choose all of the variables

Instead to create a "punishment" for including noise terms we define the adjusted R^2 as:

$$R_{adj}^2 = 1 - \frac{RSS/(n - d - 1)}{TSS/(n - 1)}$$

Adjusted R^2

Used regsubset function and specified backward method to implement step wise implementation

```
bac.sub= regsubsets(exret~., data= data_train, nvmax=10, method= "backward")
t(summary(bac.sub)$which)
```

	1	2	3	4	5	6	7	8
## (Intercept)	TRUE	TRUE	TRUE	TRUE	TRUE	TRUE	TRUE	TRUE
## alpha	FALSE	TRUE	TRUE	TRUE	TRUE	TRUE	TRUE	TRUE
## b_mkt	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	TRUE	TRUE
## b_smb	FALSE	FALSE	FALSE	TRUE	TRUE	TRUE	TRUE	TRUE
## b_hml	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	TRUE
## ivol	FALSE	FALSE	FALSE	FALSE	FALSE	TRUE	TRUE	TRUE
## tvol	FALSE	FALSE	FALSE	FALSE	TRUE	TRUE	TRUE	TRUE
## exret1	TRUE	TRUE	TRUE	TRUE	TRUE	TRUE	TRUE	TRUE
## exret_sq	FALSE	FALSE	TRUE	TRUE	TRUE	TRUE	TRUE	TRUE

- Determined the model by looking at the maximum r-squared value across different models which is model 8 using all variables

Determining the model by looking at the maximum r-squared value
summary(bac.sub)\$rsq

```
## [1] 0.007772254 0.014351234 0.016415515 0.016577997 0.016829645
0.017233883
## [7] 0.018357064 0.018358528
```

```
which.max(summary(bac.sub)$rsq)
```

```
## [1] 8
```

- Then our best model according to R-squared value is using all the variables and coefficient as follows

```
" (Intercept) alpha b_mkt b_smb b_hml ivol tvol exret1 -3.571634e-03 2.504058e+00
4.890125e-03 5.988364e-03 -9.127282e-05 8.152261e-01 -6.265090e-01 -9.563937e-02
exret_sq -1.297810e+00 "
```

Then our best model according to R-squared value is using all the variables and coefficient as follows

```
"
(Intercept)      alpha      b_mkt      b_smb      b_hml
ivol      tvol      exret1
-3.571634e-03  2.504058e+00  4.890125e-03  5.988364e-03  -9.127282e-05
8.152261e-01  -6.265090e-01  -9.563937e-02
      exret_sq
-1.297810e+00
"
```

```
## [1] "\n(Intercept)      alpha      b_mkt      b_smb      b_hml
ivol      tvol      exret1 \n-3.571634e-03  2.504058e+00  4.890125e-03
5.988364e-03  -9.127282e-05  8.152261e-01  -6.265090e-01  -9.563937e-02 \n
exret_sq \n-1.297810e+00 \n"
```

```
rsq.model= lm(exret~., data=data_train)
```

```
#summary(rsq.model)
```

```
coef(rsq.model)
```

```
## (Intercept)      alpha      b_mkt      b_smb      b_hml
## -3.571634e-03  2.504058e+00  4.890125e-03  5.988364e-03  -9.127282e-05
##      ivol      tvol      exret1      exret_sq
## 8.152261e-01  -6.265090e-01  -9.563937e-02  -1.297810e+00
```

6.a. Using the training sample, fit a Ridge regression model with all the variables to forecast excess return. Create a graph with the different values of lambda and the coefficients. Using the `cv.glmnet` function in the `glmnet` library, fit a Ridge regression model with a 10-fold cross-validation to choose the tuning parameter lambda. Print the value of the coefficients. Using the best lambda and the test sample, predict next day excess return and calculate the mean squared error.

Ridge Regression



Previously defined we have the least squares fitting procedure that estimates $\beta_0, \beta_1, \dots, \beta_p$ and minimizes

$$RSS = \sum_{i=1}^n \left(y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij} \right)^2$$

Ridge Regression is similar to least squares, but it finds the estimates that minimize:

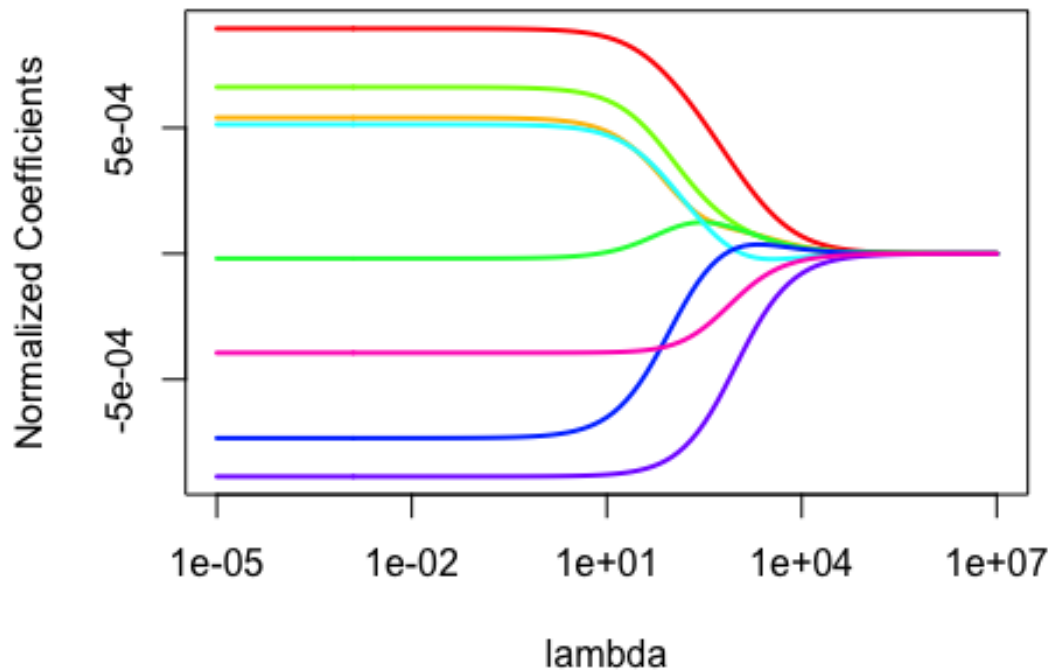
$$\sum_{i=1}^n \left(y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij} \right)^2 + \lambda \sum_{j=1}^p \beta_j^2 = RSS + \lambda \sum_{j=1}^p \beta_j^2$$

where $\lambda \geq 0$ is a *tuning parameter*

Ridge Regression

```
# Fitting a ridge regression model using training dataset
library(MASS)
grid= 10^seq(7,-5, by= -.1)
a= lm.ridge(exret~., data=data_train, lambda = grid)

# Creating a graph with the diferent values of lambda and the coefficients
leg.col= rainbow(dim(a$coef)[1])
matplot(grid,t(a$coef), type="l",lty=1,lwd=2, log = "x",col=leg.col,
xlab="lambda", ylab="Normalized Coefficients")
```



- Found the Tuning parameter which is -4.60517
- Calculated mean squared error which is 8.489139e-05

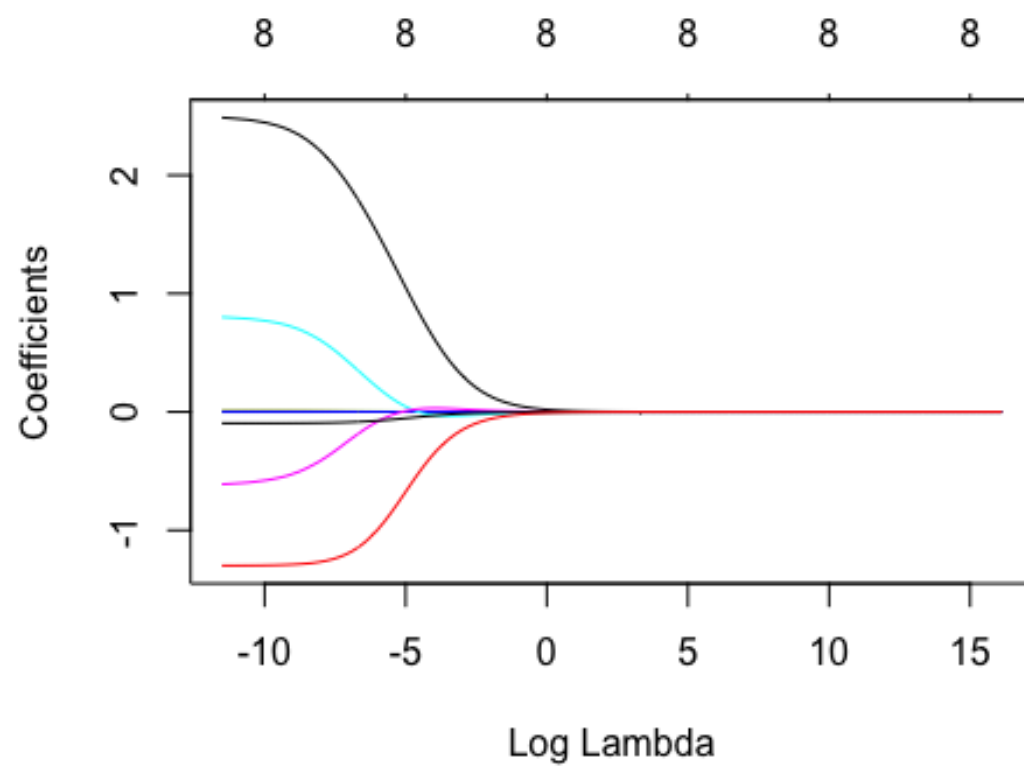
```
library(glmnet)
x= model.matrix(exret~., data=data_train)[,-1]
y= data_train$exret
ridge.mod= glmnet(x,y,alpha = 0,lambda=grid)

# 10-fold cross-validation to choose the tuning parameter Lambda
cv.out= cv.glmnet(x,y,alpha=0,lambda=grid)
bestlam= cv.out$lambda.min

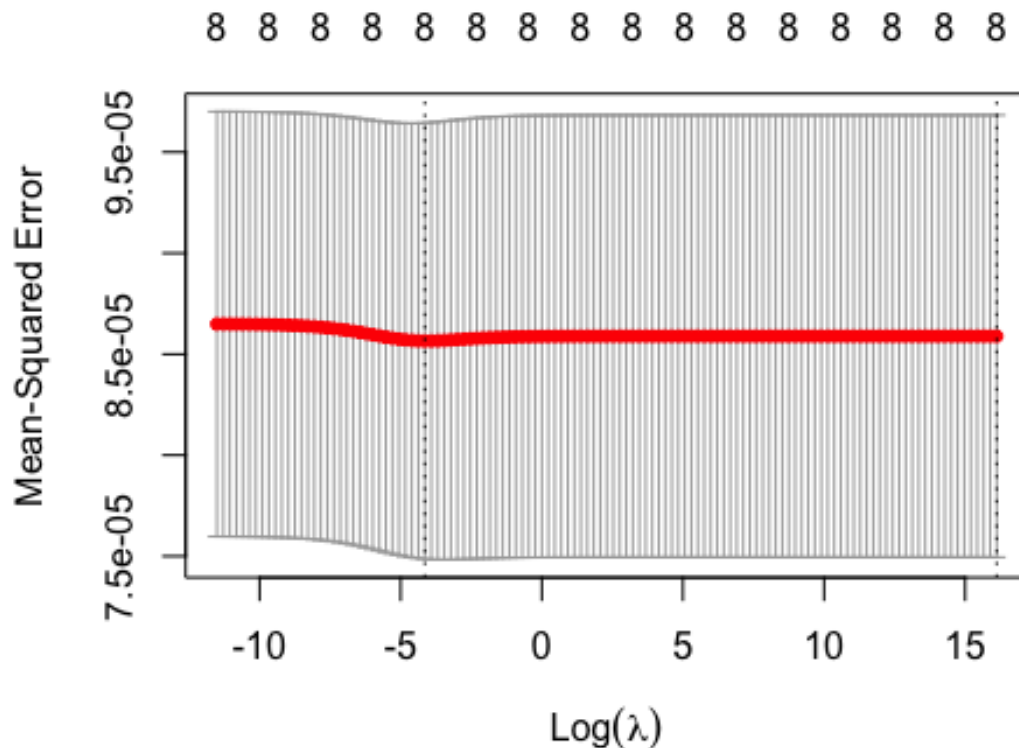
# Finding the Tuning parameter which is -4.60517
log(bestlam)

## [1] -4.144653

plot(cv.out$glmnet.fit,"lambda")
```



```
plot(cv.out)
```



```
# Predicting values using ridge regression
bestrid.pred= predict(ridge.mod, s= bestlam,newx = x)

# Calculating mean squared error which is 8.489139e-05
mse_rid= mean((bestrid.pred- y)^2)
mse_rid

## [1] 8.489139e-05
```

- Predicted Coefficients

```
# Predicting Coefficients
ridgecoef= glmnet(x,y,alpha=0,lambda = grid)
predict(ridgecoef,type="coefficients", s= bestlam)

## 9 x 1 sparse Matrix of class "dgCMatrix"
##              s1
## (Intercept) -0.0005307564
## alpha       0.6784633576
## b_mkt       0.0007076546
## b_smb       0.0007150663
## b_hml       0.0003646288
## ivol        -0.0225943656
```

```
## tvol      0.0305012556
## exret1    -0.0335433538
## exret_sq  -0.3926163992
```

b. Fit a Lasso regression model with a 10-fold cross-validation to choose the tuning parameter lambda. Print the value of the coefficients. Using the best lambda and the test sample, predict next day excess return and calculate the mean squared error.

The Lasso



A relatively recent alternative to ridge regression that picks the coefficients $\hat{\beta}_{\lambda}^L$ that minimizes:

$$\sum_{i=1}^n \left(y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij} \right)^2 + \lambda \sum_{j=1}^p |\beta_j| = RSS + \lambda \sum_{j=1}^p |\beta_j|$$

This approach produces **sparse** models.

Lasso Regression

- Calculated tuning parameter which is -8.289306
- Calculated mse value which is 8.457083e-05
- Predicted Coefficients

Fitting a Lasso model

```
lasso.mod= glmnet(x,y,alpha=1,lambda=grid)
```

10-fold cross-validation to choose the tuning parameter Lambda

```
cv.out=cv.glmnet(x,y,alpha=1, lambda=grid)
```

#plot(cv.out)

```
bestlam=cv.out$lambda.min
```

```
log(bestlam)
```

```
## [1] -8.289306
```



```

# predicting excess return
bestlas.pred= predict(lasso.mod, s=bestlam, newx=x)

# Predicting Coefficients
lassocoeff= glmnet(x,y,alpha=1,lambda = grid)
predict(lassocoeff,type="coefficients", s= bestlam)

## 9 x 1 sparse Matrix of class "dgCMatrix"
##              s1
## (Intercept) -8.263723e-05
## alpha      1.401691e+00
## b_mkt      .
## b_smb      .
## b_hml      .
## ivol       .
## tvol       .
## exret1     -6.274945e-02
## exret_sq   -4.088129e-01

# mse using lasso, ,mse value is 8.457083e-05
mean((bestlas.pred-y)^2)

## [1] 8.457083e-05

```

c. Compare and discuss the results of Lasso and Ridge regression indicating what approach you will choose and why.

- Both Lasso and Ridge Regression gives very small Mean Squared Error, respectively 8.457083e-05, 8.471404e-05.
- Lasso uses -8.059048 as tuning parameter, whereas ridge uses -4.60517 as tuning parameters
- It's wise to say that lasso is better in this case, since it shrinks the coefficients exactly to zero which gives the model a good interpretation and makes it more simple

*# Both Lasso and Ridge Regression gives very small Mean Squared Error, respectively 8.457083e-05, 8.471404e-05.
Lasso uses -8.059048 as tuning parameter, whereas ridge uses -4.60517 as tuning parameters*

it's wise to say that lasso is better in this case, since it shrinks to coefficients exactly to zero which gives the model a good interpretation and makes it more simple

Question 2

Create another field “Direction” in this data frame that looks to the direction of the excess return of the next period (exret1), this direction should be listed as a factor, not a number. After “Direction” is created, exret1 should not be included in the dataset.

1. Using the training set, run LDA to forecast “Direction” on the training set. Predict with the test sample. Calculate the confusion matrix and accuracy.

```
# Adding Direction to the data, and assigning value 1 if exret1 >= 0 or 0 if exret1 <= 0

data$Direction= 0

for (i in seq(1: nrow(data) ) ) {

  if (data$exret1[i]>=0){

    data$Direction[i]=1

  } else {

    data$Direction[i]=0

  }

}

# Setting direction as factor
data$Direction= as.factor(data$Direction)

table(data$Direction)

##
##    0    1
## 734 775

# Removing exret1 after Direction variable is created
data= subset(data, select= -exret1)
```

```

#head(data)

# Updating test and training data as well
data_train= data[1: train, ]
data_test= data[(train+1): nrow(data), ]

head(data_test)

##      alpha  b_mkt   b_smb   b_hml      ivol      tvol      exret
exret_sq
## 1057 -6e-04 0.8333 -0.1090 -0.2647 0.008877 0.010471 -0.010955 2.152089e-
06
## 1058 -6e-04 0.8318 -0.1118 -0.2631 0.008876 0.010457 -0.001467 1.032662e-
04
## 1059 -6e-04 0.8279 -0.1073 -0.2591 0.008843 0.010413 -0.010162 2.033476e-
06
## 1060 -6e-04 0.8303 -0.1156 -0.2565 0.008835 0.010448 -0.001426 2.724840e-
05
## 1061 -6e-04 0.8275 -0.1149 -0.2650 0.008821 0.010431 -0.005220 1.823290e-
05
## 1062 -6e-04 0.8341 -0.0948 -0.2560 0.008818 0.010429 0.004270 5.712100e-
06
##      Direction
## 1057         0
## 1058         0
## 1059         0
## 1060         0
## 1061         1
## 1062         1

```

- Created the confusion matrix
- Calculated accuracy which is 0.5055188

```

# Fitting LDA for Direction on training dataset
lda.fit= lda(Direction~.,data=data_train )
#lda.fit

# Predicting the direction using train data
lda.pred_train= predict(lda.fit, data_train)
lda.class_train= lda.pred_train$class

# Confusing matrix using train Data
#table(lda.class_train, data_train$Direction)
# Accuracy using training data
train_accur= mean(lda.class_train== data_train$Direction)

```

```

# Predicting the direction with test sample
lda.pred= predict(lda.fit,data_test)
lda.class= lda.pred$class

# Creating the Confusion Matrix
table(lda.class,data_test$Direction)

##
## lda.class    0    1
##           0 133 148
##           1   76  96

# Calculating accuracy
test_accur= mean(lda.class== data_test$Direction)

#(133+ 96) / (96+133+148 +76)

## Model has higher accuracy using training data in comparison to test data(
0.5055188 )

Accuracy_table= list("Train Accuracy " = train_accur, "Test
Accuracy"=test_accur )

data.frame(Accuracy_table)

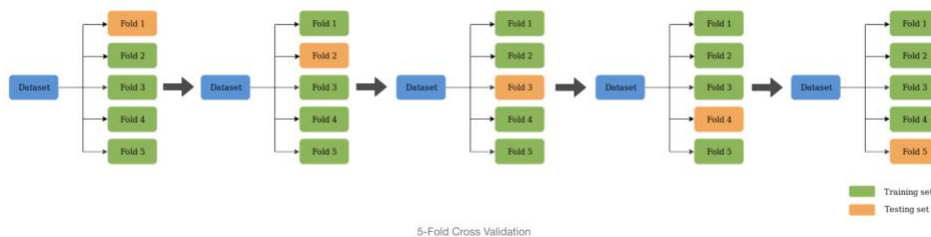
##   Train.Accuracy. Test.Accuracy
## 1           0.5397727           0.5055188

```

2. Create code to determine the estimate of the expected test error of your model to forecast “Direction” using K=5 cross validation. Do this by actually splitting the complete dataset into five pieces (can use function cut) and give the average of the test error, not just by using a command from a package.

What is K-Fold Cross Validation?

K-Fold CV is where a given data set is split into a K number of sections/folds where each fold is used as a testing set at some point. Lets take the scenario of 5-Fold cross validation ($K=5$). Here, the data set is split into 5 folds. In the first iteration, the first fold is used to test the model and the rest are used to train the model. In the second iteration, 2nd fold is used as the testing set while the rest serve as the training set. This process is repeated until each fold of the 5 folds have been used as the testing set.



K Fold Cross Validation

- Data split into 5 equal pieces

K Fold Cross Validation, used boot library

Splitting data into 5 equal pieces

```
data_1= data[(1: (nrow(data)/5)), ]
```

```
#tail(data_1)
```

```
data_2= data[(nrow(data_1)+1): (2*(nrow(data)/5)) , ]
```

```
#tail(data_2)
```

```
data_3= data[(2*(nrow(data)/5)+1): (3*(nrow(data)/5)) , ]
```

```
#tail(data_3)
```

```
data_4= data[(3*(nrow(data)/5)): (4*(nrow(data)/5)) , ]
```

```
#tail(data_4)
```

```
data_5= data[(4*(nrow(data)/5)):nrow(data) , ]
```

```
tail(data_5)
```

```
##      alpha  b_mkt  b_smb  b_hml    ivol    tvol    exret    exret_sq
## 1503 1e-04  0.6321 -0.4551 -0.3169 0.009103 0.011077 -0.006939 1.034909e-05
## 1504 1e-04  0.6291 -0.4540 -0.3131 0.009103 0.011031 -0.003217 4.019560e-07
## 1505 1e-04  0.6280 -0.4545 -0.3124 0.009102 0.010987  0.000634 1.874890e-07
## 1506 2e-04  0.6043 -0.4185 -0.3261 0.009063 0.010791  0.000433 1.195742e-04
## 1507 2e-04  0.6032 -0.4074 -0.3288 0.009090 0.010608 -0.010935 2.524058e-05
```

```
## 1508 2e-04 0.5978 -0.3991 -0.3294 0.009079 0.010568 0.005024 5.230369e-06
##      Direction
## 1503      0
## 1504      1
## 1505      1
## 1506      0
## 1507      1
## 1508      1
```

- The average of the test error is 0.505309 using using k fold cross validation where k=5

data_1 is the training set

```
lda.data_1= lda(Direction~., data= rbind(data_2,data_3,data_4,data_5) )
data_1.pred= predict(lda.data_1, newdata = data_1)
err_1= mean(data_1.pred$class!=data_1$Direction)
```

data_2 is the test set

```
lda.data_2= lda(Direction~., data= rbind(data_1,data_3,data_4,data_5) )
data_2.pred= predict(lda.data_2, newdata = data_2)
err_2= mean(data_2.pred$class!=data_2$Direction)
```

data_3 is the test set

```
lda.data_3= lda(Direction~., data= rbind(data_1,data_2,data_4,data_5) )
data_3.pred= predict(lda.data_3, newdata = data_3)
err_3= mean(data_3.pred$class!=data_3$Direction)
```

data_4 is the test set

```
lda.data_4= lda(Direction~., data= rbind(data_1,data_2,data_3,data_5) )
data_4.pred= predict(lda.data_4, newdata = data_4)
err_4= mean(data_4.pred$class!=data_4$Direction)
```

data_5 is the test set

```
lda.data_5= lda(Direction~., data= rbind(data_1,data_2,data_3,data_4) )
data_5.pred= predict(lda.data_5, newdata = data_5)
err_5= mean(data_5.pred$class!=data_5$Direction)
```

```
err_list= cbind(err_1,err_2,err_3,err_4,err_5)
```

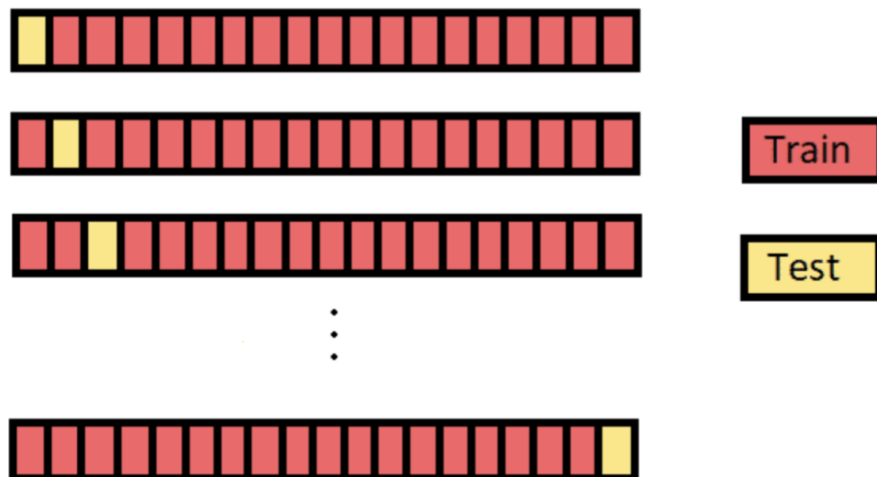
```
data.frame(err_list)

##      err_1      err_2      err_3 err_4      err_5
## 1 0.5149502 0.4768212 0.5016611  0.5 0.5331126

# The average of the test error 0.505309
mean(err_list)

## [1] 0.505309
```

3. Determine the LOOCV estimate of the expected test error of your model to forecast “Direction” using the complete dataset. How do your answers to each part of this question compare? Do you see any noticeable differences between your answers? Why do you think that is?



Leave One Out Cross validation. This is another cross validation

K Fold Cross Validation

Mean error is by using LOOCV is 0.4910537 whereas mean error using K-fold is 0.505309 they are so close to each other

Obviously LOOCV takes more time than K-Fold to come up with MSE, since the number of folds equals to the number of observation -1

LOOCV can be subject to high variance or overfitting, as we are feeding the model almost all the training data to learn and just a single observation to evaluate.

K-Fold reduces the variance shown by LOOCV and introduces some bias by holding out a substantially large validation set.

```
err_cont=c()
#data
for (i in 1:nrow(data)) {
  train <- data[-i,]
  test <- data[i,]

  lda.fit= lda(Direction~., data= train)

  lda.pred= predict(lda.fit, test)

  lda.class=lda.pred$class

  err= mean( lda.class!= test$Direction )

  err_cont= c(err_cont,err)

}

mean(err_cont)
## [1] 0.4910537

# Mean error is by using LOOCV is 0.4910537 whereas mean error using K-fold
is 0.505309 they are so close to each other

# Obviously LOOCV takes more time than K-Fold to come up with MSE
```


LOOCV can be subject to high variance or overfitting

K-Fold reduces the variance shown by LOOCV and introduces some bias by holding out a substantially large validation set.

Question 3

This question should be answered using the Weekly data set, which is part of the ISLR package. This data contains 1,089 weekly returns for 21 years, from the beginning of 1990 to the end of 2010.

1. What does the data represent?

- Data represents weekly return with lag values till 5 and Volume, and Direction
- Min, Median and Quartiles values of the data represented in below

Data represents weekly return with lag values till 5 and Volume, and Direction

Min, Median and Quartiles values of the data represented in below

`library(ISLR)`

`head(Weekly)`

```
##   Year  Lag1  Lag2  Lag3  Lag4  Lag5  Volume  Today Direction
## 1 1990  0.816  1.572 -3.936 -0.229 -3.484 0.1549760 -0.270      Down
## 2 1990 -0.270  0.816  1.572 -3.936 -0.229 0.1485740 -2.576      Down
## 3 1990 -2.576 -0.270  0.816  1.572 -3.936 0.1598375  3.514       Up
## 4 1990  3.514 -2.576 -0.270  0.816  1.572 0.1616300  0.712       Up
## 5 1990  0.712  3.514 -2.576 -0.270  0.816 0.1537280  1.178       Up
## 6 1990  1.178  0.712  3.514 -2.576 -0.270 0.1544440 -1.372      Down
```

`summary(Weekly)`

```
##      Year      Lag1      Lag2      Lag3
## Min.   :1990  Min.   :-18.1950  Min.   :-18.1950  Min.   :-18.1950
## 1st Qu.:1995  1st Qu.: -1.1540  1st Qu.: -1.1540  1st Qu.: -1.1580
## Median :2000  Median :  0.2410  Median :  0.2410  Median :  0.2410
## Mean   :2000  Mean   :  0.1506  Mean   :  0.1511  Mean   :  0.1472
## 3rd Qu.:2005  3rd Qu.:  1.4050  3rd Qu.:  1.4090  3rd Qu.:  1.4090
## Max.   :2010  Max.   : 12.0260  Max.   : 12.0260  Max.   : 12.0260
##      Lag4      Lag5      Volume      Today
## Min.   :-18.1950  Min.   :-18.1950  Min.   :0.08747  Min.   :-18.1950
## 1st Qu.: -1.1580  1st Qu.: -1.1660  1st Qu.:0.33202  1st Qu.: -1.1540
## Median :  0.2380  Median :  0.2340  Median :1.00268  Median :  0.2410
## Mean   :  0.1458  Mean   :  0.1399  Mean   :1.57462  Mean   :  0.1499
## 3rd Qu.:  1.4090  3rd Qu.:  1.4050  3rd Qu.:2.05373  3rd Qu.:  1.4050
## Max.   : 12.0260  Max.   : 12.0260  Max.   :9.32821  Max.   : 12.0260
## Direction
## Down:484
## Up  :605
```

```
##  
##  
##  
##
```

2. Use the full data set to perform a logistic regression with Direction as the response and the five lag variables plus Volume as predictors. Use the summary function to print the results. Do any of the predictors appear to be statistically significant? If so, which ones?

- According to summary results, Yes some of the variables statistically significant.
- Intercept and Lag value 2 values are significant, it could be also said that Lag 1(p-value 0.1181) is somewhat significant

```
# According to summary results, Yes some of the variables statistically  
significant.  
# Intercept and Lag value 2 values are significant, it could be also said  
that Lag 1 is somewhat significant
```

```
log.fit= glm(Direction~ Lag1+ Lag2 +Lag3 +Lag4 +Lag5+Volume, data=Weekly,  
family = binomial)  
summary(log.fit)
```

```
##  
## Call:  
## glm(formula = Direction ~ Lag1 + Lag2 + Lag3 + Lag4 + Lag5 +  
##      Volume, family = binomial, data = Weekly)  
##  
## Deviance Residuals:  
##      Min       1Q   Median       3Q      Max   
## -1.6949  -1.2565   0.9913   1.0849   1.4579   
##  
## Coefficients:  
##              Estimate Std. Error z value Pr(>|z|)      
## (Intercept)  0.26686    0.08593   3.106   0.0019 **      
## Lag1        -0.04127    0.02641  -1.563   0.1181        
## Lag2         0.05844    0.02686   2.175   0.0296 *       
## Lag3        -0.01606    0.02666  -0.602   0.5469        
## Lag4        -0.02779    0.02646  -1.050   0.2937        
## Lag5        -0.01447    0.02638  -0.549   0.5833        
## Volume       -0.02274    0.03690  -0.616   0.5377        
## ---  
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1  
##  
## (Dispersion parameter for binomial family taken to be 1)  
##  
##      Null deviance: 1496.2  on 1088  degrees of freedom  
## Residual deviance: 1486.4  on 1082  degrees of freedom  
## AIC: 1500.4
```

```
##
## Number of Fisher Scoring iterations: 4

coef(log.fit)

## (Intercept)      Lag1      Lag2      Lag3      Lag4      Lag5
##  0.26686414 -0.04126894  0.05844168 -0.01606114 -0.02779021 -0.01447206
##      Volume
## -0.02274153

# Predicting prob with given logistic regression model. type = "response"
# option tells R to output probabilities of the form  $P(Y = 1|X)$ 

log.probs=predict(log.fit,type="response")

log.probs[1:10]

##      1      2      3      4      5      6      7
## 0.6086249 0.6010314 0.5875699 0.4816416 0.6169013 0.5684190 0.5786097
## 0.5151972
##      9     10
## 0.5715200 0.5554287
```

3. Fit a logistic regression model using a training data period from 1990 to 2008, using the predictors from the previous problem that you determined were statistically significant. Test your model on the held out data (that is, the data from 2009 and 2010) and express its accuracy.

- Accuracy is calculated which is 0.5769231

```
# Training data creation
Weekly_train= Weekly[Weekly["Year"] < 2009,]
#tail(Weekly_train)

# Test data creation
Weekly_test= Weekly[Weekly["Year"] >= 2009,]
#head(Weekly_test)
# Fitting a logistic regression with predictors found significant in the
# model above, and used only training data
log.fit_train= glm(Direction~Lag2+Lag1, data=Weekly_train, family=binomial)

summary(log.fit_train)

##
## Call:
## glm(formula = Direction ~ Lag2 + Lag1, family = binomial, data =
## Weekly_train)
##
```

```

## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.6149  -1.2565   0.9989   1.0875   1.5330
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)   0.21109    0.06456   3.269  0.00108 **
## Lag2          0.05384    0.02905   1.854  0.06379 .
## Lag1         -0.05421    0.02886  -1.878  0.06034 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 1354.7  on 984  degrees of freedom
## Residual deviance: 1347.0  on 982  degrees of freedom
## AIC: 1353
##
## Number of Fisher Scoring iterations: 4

# Predicting values using training data

weekly_prob_train= predict(log.fit_train,type="response")

glm_pred_train= rep(0,985)
glm_pred_train[weekly_prob_train>0.5]=1

# Assigning 0 and 1 for the Direction Down and Up
levels(Weekly_train$Direction)= c(0,1)

# Creating Confusion Matrix for the training data
#table(glm_pred_train,Weekly_train$Direction)

# Calculating the accuracy
#mean(glm_pred_train== Weekly_train$Direction)

# Testing the model on testing data set
weekly_prob= predict(log.fit_train,newdata = Weekly_test, type= "response")
glm_pred= rep(0,104)
glm_pred[weekly_prob>0.5]=1

# Confusion Matrix
table(glm_pred,Weekly_test$Direction )

##
## glm_pred Down Up

```

```
##           0    7    8
##           1   36  53

# Changing the Level for Direction, 0 for Down, 1 for Up
levels(Weekly_test$Direction)= c(0,1)

# Confusion Matrix modified with only 0 and 1s
table(glm_pred,Weekly_test$Direction )

##
## glm_pred  0  1
##           0  7  8
##           1 36 53

# Calculating the Accuracy
## Accuracy is 0.5769231
mean(glm_pred== Weekly_test$Direction)

## [1] 0.5769231
```

4. Repeat Part 3 using LDA.

- Calculated Accuracy which is 0.5769231

```
library(MASS)

# Fitting a Model Using Linear Discriminant Analysis
lda.fit= lda(Direction~Lag2+Lag1, data= Weekly_train)
lda.predict= predict(lda.fit, Weekly_test)

lda.class= lda.predict$class

# Creating a confusion matrix
table(lda.class,Weekly_test$Direction )

##
## lda.class  0  1
##           0  7  8
##           1 36 53

# Calculating the Accuracy
## Accuracy is 0.5769231
mean(lda.class== Weekly_test$Direction)

## [1] 0.5769231
```

5. Repeat Part 3 using QDA.

- Calculated Accuracy which is 0.5576923

```
# Fitting a model using Quadratic Discriminant Analysis
qda.fit= qda(Direction~Lag2+Lag1, data= Weekly_train)
```

```

# Prediction Direction with QDA
qda.predict= predict(qda.fit,Weekly_test)
qda.class= qda.predict$class

# Creating a confusion matrix
table(qda.class, Weekly_test$Direction)

##
## qda.class  0  1
##          0  7 10
##          1 36 51

# Calculating the Accuracy
## Accuracy is 0.5576923
mean(qda.class== Weekly_test$Direction)

## [1] 0.5576923

```

6. Repeat Part 3 using KNN with K = 1, 2, 3.

- K=1, Calculated the Accuracy which is 0.4807692
- K=2, Calculated the Accuracy which is 0.4615385
- K=3, Calculated the Accuracy which is 0.5192308

```

library(class)

## Fitting a model using K- Nearest Neighbor using k=1

attach(Weekly_train)
train.X= cbind(Lag1,Lag2)
train.Direction= Weekly_train$Direction
attach(Weekly_test)
test.X= cbind(Lag1,Lag2)

knn.pred= knn(train=train.X,test=test.X, cl=train.Direction, k=1)

# Calculating the Accuracy
# Accuracy is 0.4807692
mean(knn.pred==Weekly_test$Direction)

## [1] 0.4807692

## Fitting a model using K- Nearest Neighbor using k=2
# Calculating the Accuracy
# Accuracy is 0.4615385

```

```

knn.pred= knn(train=train.X,test=test.X, cl=train.Direction, k=2)
mean(knn.pred==Weekly_test$Direction)

## [1] 0.5480769

## Fitting a model using K- Nearest Neighbor using k=3
knn.pred= knn(train=train.X,test=test.X, cl=train.Direction, k=3)
# Calculating the Accuracy
# Accuracy is 0.5192308
mean(knn.pred==Weekly_test$Direction)

## [1] 0.5192308

```

7. Which of these methods in Parts 3, 4, 5, and 6 appears to provide the best results on this data?

- Logistic Regression Model gives Accuracy= 0.5769231
- Linear Discriminant Analysis also gives Accuracy= 0.5769231
- Quadratic Discriminant Analysis gives Accuracy= 0.5576923
- KNN with K=1 gives Accuracy=0.4807692
- KNN with K=2 gives Accuracy=0.4615385
- KNN with K=3 gives Accuracy=0.5192308
- Among all of these different techniques, Logistic Regression and Linear Discriminant Analysis gave the highest Accuracy

```

# Logistic Regression Model gives Accuracy= 0.5769231
# Linear Discriminant Analysis also gives Accuracy= 0.5769231
# Quadratic Discriminant Analysis gives Accuracy= 0.5576923
# KNN with K=1 gives Accuracy=0.4807692
# KNN with K=2 gives Accuracy=0.4615385
# KNN with K=3 gives Accuracy=0.5192308

## Among all of these different techniques, Logistic Regression and Linear
Discriminant Analysis gave the highest Accuracy

```

Question 4

Write a function that works in R to gives you the parameters from a linear regression on a data set of n predictors. You can assume all the predictors and the prediction is numeric. Include in the output the standard error of your variables. You cannot use the `lm` command in this function or any of the other built in regression models.



Denote by \mathbf{X} the $N \times (p + 1)$ matrix with each row an input vector with a 1 in the first column and let \mathbf{y} be the N vector of outputs. Using these definitions we have:

$$RSS(\beta) = (\mathbf{y} - \mathbf{X}\beta)^T(\mathbf{y} - \mathbf{X}\beta)$$

which is a quadratic function of the $p + 1$ parameters.

÷

Differentiating with respect to β gives us:

$$\begin{aligned}\frac{\partial RSS}{\partial \beta} &= -2\mathbf{X}^T(\mathbf{y} - \mathbf{X}\beta) \\ \frac{\partial^2 RSS}{\partial \beta \partial \beta^T} &= 2\mathbf{X}^T\mathbf{X}\end{aligned}$$

Least Square Method-1



If we assume that \mathbf{X} has full column rank, and so $\mathbf{X}^T \mathbf{X}$ is positive definite, then we set the first derivative to 0:

$$\mathbf{X}^T (\mathbf{y} - \mathbf{X}\beta) = 0$$

which gives us the solution:

$$\hat{\beta} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

and so the fitted values are:

$$\hat{\mathbf{y}} = \mathbf{X}\hat{\beta} = \mathbf{X}(\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

Least Square Method-2

$$\mathbf{X}'\mathbf{X} = \begin{bmatrix} n & \sum_{i=1}^n x_i \\ \sum_{i=1}^n x_i & \sum_{i=1}^n x_i^2 \end{bmatrix}$$

Least Square Method-3

$$\mathbf{b} = \begin{bmatrix} b_0 \\ b_1 \\ \vdots \\ b_{p-1} \end{bmatrix} = (\mathbf{X}'\mathbf{X})^{-1} \mathbf{X}'\mathbf{Y}$$

Least Square Method-4

- Created a linear regression function, note column number should be specified to select the dependent variable

Creating a linear regression function, note column number should be specified to select the dependent variable

```
regress= function(dataset,col_num){  
  
  Y= as.matrix(dataset[col_num])  
  X= as.matrix (dataset[-col_num])  
  
  # Adding 1s to the Matrix  
  X = cbind(rep(1,nrow(dataset)), X )  
}
```

```

a= t(X) %*% X
b= t(X) %*% Y
a_inv= solve(a)
beta= a_inv %*% b

colnames(beta)= "coefficients"

N=nrow(X)
p= ncol(X)
df= N-p-1

# Calculating Standard Error of the coefficients
sig_sq= 1/(N-p) * sum((Y- X%*%beta ) ^2 )
SE= diag(sig_sq *a_inv)
SE= sqrt(SE)

# Calculating t-value

result= cbind(beta, SE)

result= data.frame(result)
result$t_value= result$coefficients/result$SE

# Calculating p- value

#2*pt(q=12.5785546,df= 5,lower.tail = F)

#sapply(result$t_value, pt)

return(result)

}

```

- Testing the function on a dataset
- Both lm function and own built function gives the same result

```

library(readr)
# Testing the regression function using data set from the URL below

data_set <-
read_delim("https://online.stat.psu.edu/stat462/sites/onlinecourses.science.psu.edu.stat462/files/data/soapsuds/index.txt", "\t", escape_double =
FALSE, trim_ws = TRUE)

```

```
# Adding another variable to the dataset to make it multivariable data set  
data_set["soap_sq"] = data_set$soap^2  
#head(data_set)
```

```
# Using lm function to create a regression model  
lm(suds~soap+soap_sq,data=data_set)
```

```
##  
## Call:  
## lm(formula = suds ~ soap + soap_sq, data = data_set)  
##  
## Coefficients:  
## (Intercept)      soap      soap_sq  
##      -34.714      21.548      -1.095
```

```
# Using regress own built function to create a regression model, both gives  
the same result  
regress(data_set,col_num=2)
```

```
##      coefficients      SE  t_value  
##      -34.714286 23.9175170 -1.451417  
## soap      21.547619  8.9013553  2.420712  
## soap_sq   -1.095238  0.8067178 -1.357647
```